

# Parallel Programming SS21 Final Project

02 Gaussian Elimination

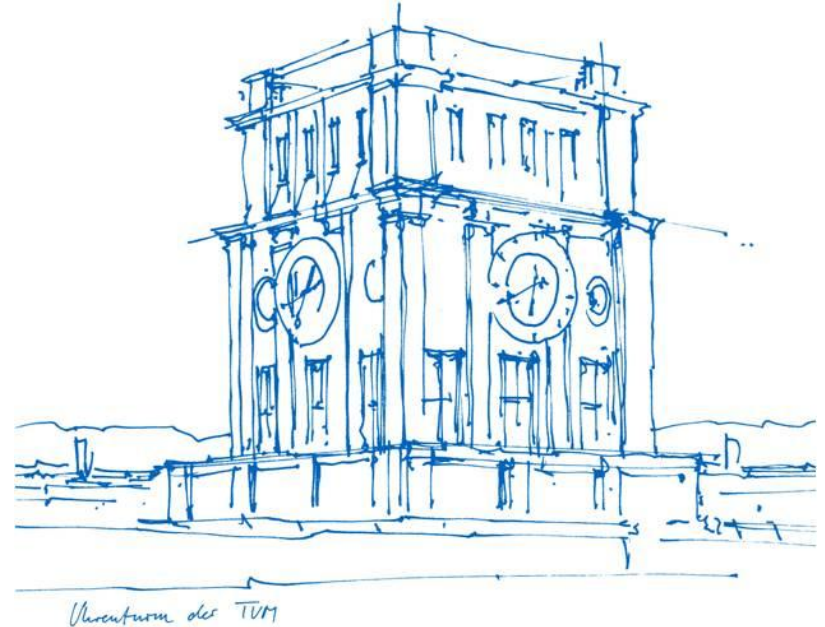
Group 218

13.07.2021

Sandro Bauer

Marcel Braasch

Syed Salman Ahmed



# Outline



1. Sequential code analysis
2. OpenMP
3. MPI
4. Hybrid
5. Bonus
6. Conclusion

# Sequential code analysis - Forward Elimination

When calculating the forward elimination, for every row, we take the diagonal element  $d$ , calculate an elimination factor  $f=d/x$  such that for all lower elements in the column  $r \cdot x - r = 0$ . This diagonal element stays constant with respect to  $x$ . Thus we do not need to calculate it over and over again.

Graphically, that is


$$d_1 = 1$$

$$x_{11} = 6$$

$$f_1 = \frac{1}{6}$$

$$x_{11} = x_{11} \cdot f_1 - x_{11} = 0$$

Repeat for all  $x_{1j}$ .



1	2	3	1	a	1
6	-4	5	0	b	6
2	3	5	0	c	9
5	-10	1	0	d	0

=

1
6
9
0

⇒ We can see that  $d_1$  stays constant for all eliminations. The given algorithm constantly calculates  $d_1$ , thus we moved this calculation outside. All further analyses apply to the algorithm with this change.

# Sequential code analysis - Theoretical Analysis

For the Forward Elimination we have

$n(n+1)/2$  division

$(2n^3+3n^2-5n)/6$  multiplications

$(2n^3+3n^2-5n)/6$  subtractions

$= n(n+1)/2 + (2n^3+3n^2-5n)/3 = \mathbf{FE(n)}$  operations

For the Backward Elimination we have

$n(n-1)/4$  multiplications

$n(n-1)/4$  additions

$= n(n-1)/2 = \mathbf{BE(n)}$  operations

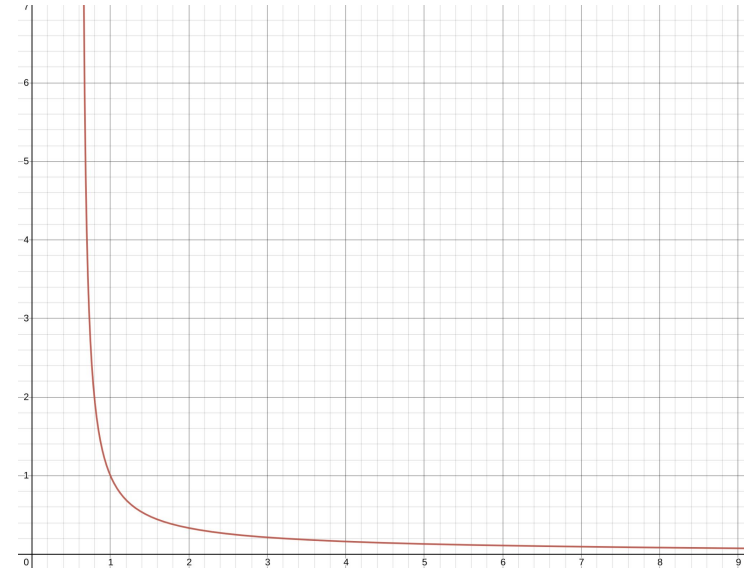
Let's look at the ratio for  $\mathbf{BE(n)/FE(n)}$ !

⇒ Quickly converges to 0

⇒ FO outweighs BO by **a lot**

⇒ E.g., for  $n = 1024$ , FO causes 99,85% of the computations

⇒ Gives rise to only optimizing FO due to high overhead



The graph  $f(n)=BE(n)/FE(n)$  showing the ratio of BE and FE indicating the importance of BE

# Sequential code analysis - Profiling and Backward



- Perf output roughly coincides with the theoretical calculation
- FE percentage slightly lower since other operations cause overhead
- Backward elimination not even visible

The backward algorithm is

1. doing a negligible amount of work
  2. has strong sequential data dependency
- ⇒ Neglect the backward algorithm throughout all our future optimizations and the rest of the presentation

Samples: 520 of event 'cache-references', Event count (approx.): 122283612

Overhead	Command	Shared Object	Symbol
97,42%	serialge	serialge	[.] Serial::ForwardEliminationOptimized
0,36%	serialge	serialge	[.] Serial::SerialSolve
0,35%	serialge	libstdc++.so.6.0.28	[.] std::num_put<char, std::ostreambuf_iterator<char, std::char_traits<char>>>
0,19%	serialge	libstdc++.so.6.0.28	[.] std::num_get<char, std::istreambuf_iterator<char, std::char_traits<char>>>
0,14%	serialge	libstdc++.so.6.0.28	[.] std::istream::sentry::sentry
0,13%	serialge	libstdc++.so.6.0.28	[.] std::num_get<char, std::istreambuf_iterator<char, std::char_traits<char>>>
0,13%	serialge	libc-2.31.so	[.] __GI___strtod_l_internal
0,08%	serialge	libstdc++.so.6.0.28	[.] std::string::reserve
0,07%	serialge	libc-2.31.so	[.] malloc
0,06%	serialge	libstdc++.so.6.0.28	[.] std::string::_Rep::_M_clone
0,05%	serialge	libc-2.31.so	[.] __strlen_avx2
0,05%	serialge	libstdc++.so.6.0.28	[.] std::__convert_to_v<double>
0,05%	serialge	[kernel.kallsyms]	[k] clear_page_erms
0,04%	serialge	serialge	[.] ReadLine

Figure: Perf profiling output for serial implementation

# Sequential code analysis - Amdahl's law

Q: What portion of the algorithm is parallelizable?



- Low data dependency
- Each of  $i+1^{\text{th}}$  rows is only dependent on the element  $A_{ii}$
- Therefore, propagating down the diagonal element can be nicely parallelized
- As this is all the algorithm is doing (besides a few other operation),  
**we expect roughly 97%** (according to Perf) to be parallelizable

# Sequential code analysis - Amdahl's law

**Q:** What portion of the algorithm is parallelizable?

So let's evaluate this, given our OMP speed up of  $L = 14$ .

Amdahl's Law is given by

$$L_p(s) = 1 / (1 - p + p/s)$$

where  $L$  is the latency (i.e., the speed up),  $p$  is the portion to be parallelizable and  $s$  the amount of workers.

Plugging in all the values yields

$$\begin{aligned} 14 &= 1 / (1 - p + p/s) \\ \Leftrightarrow p &\approx 0.96 \end{aligned}$$

⇒ According to a speed up of factor 14, approximately 96% of the program is parallelizable

Similarly, we can ask ...

# Sequential code analysis - Amdahl's law

**Q:** Given the portion of the algorithm to parallelizable ( $p=0.97$ ), what's the theoretical speed up maximum?

Plotting the  $L_p(s)$  yields

- ⇒ Assuming 97% parallelizability using 32 CPUs theoretical speed up of factor 16.5
- ⇒ Clearly **sublinear** but still very nice curve.
- ⇒ Theoretical max speed up:  
 $L_p(s)$  converges towards 33.3 for  $s \rightarrow \infty$ .

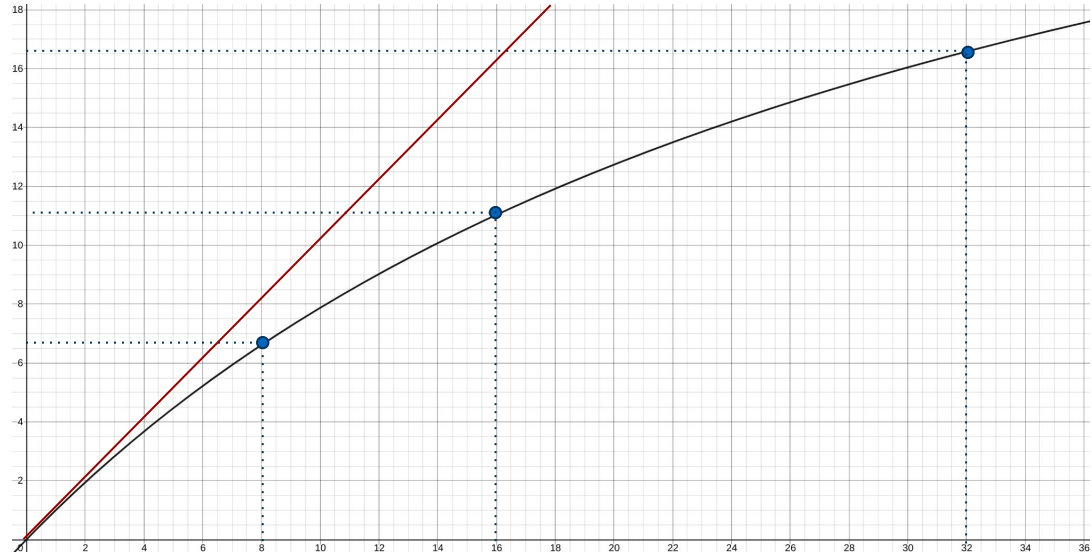


Figure: Amdahl's law with  $p=0.97$ .



# OpenMP - Parallelized implementation and approach



## Initial approach

- Used **OMP task** to parallelize program
- Resulted in **no speed up at all**
- Taking a closer look we recognized **not the correct approach**
- OMP Tasks are (loosely speaking) a way to dynamically distribute workload
- Use case would be recursion with no interdependence between tree branches
- GE has low data dependence in forward algorithm, however,
- **Dynamically distributing matrix rows very inefficient**

## Current approach

- For every pivot row from step  $i=1$ , we have to modify the following  $n-i$  rows
- All following  $n-i$  rows are only dependent on  $i$ -th row's pivot element
- Statically spread the workload across all workers minimizing overhead

# OpenMP - Parallelized implementation and approach

We found that

1. Out of the explicit scheduling schemes (static, dynamic, guided) static is the fastest
2. Runtime system using `schedule(auto)` picks static too → static seems like a good strategy

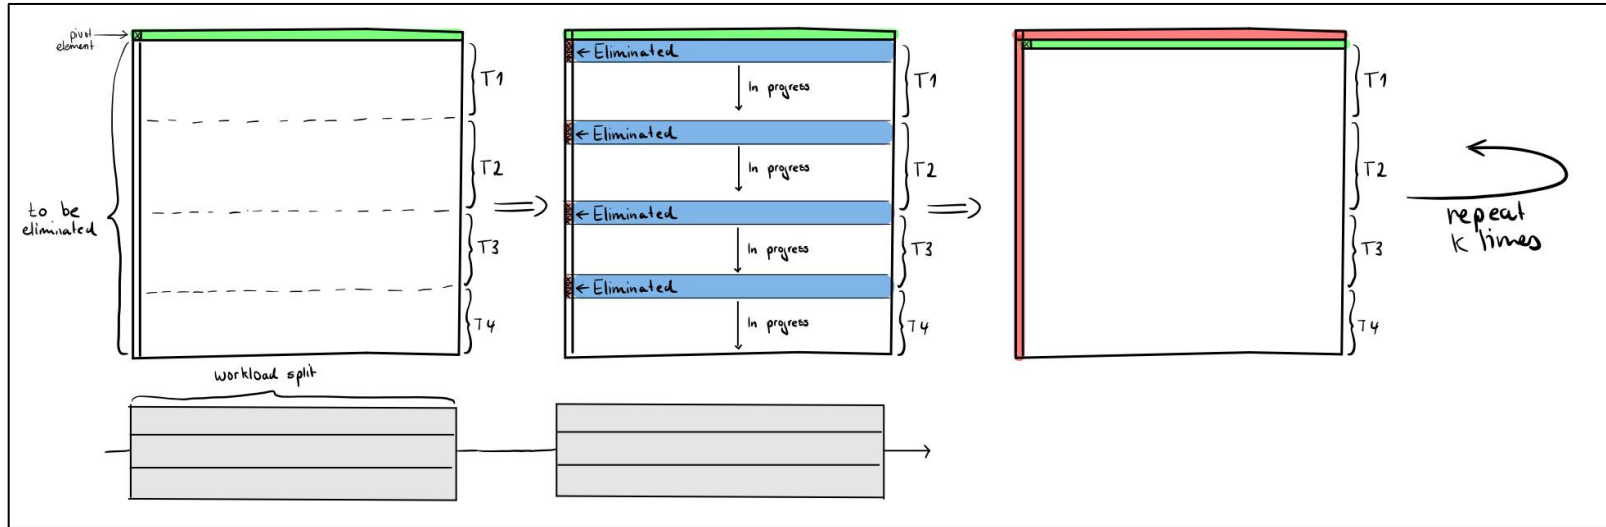


Figure: Schematic workload spread for  $k \times k$  in forward algorithm with  $n=4$  and static scheduling.

# OpenMP - Overhead and when to parallelize

**Q:** At what problem size does parallelization make sense?

- Ran every problem size 10 times and averaged over computation vs. overhead
- For small problems overhead takes almost 100% of computing time  $\Rightarrow$  parallelization makes no sense
- For the sample problems given, starting from size 512 overhead starts to decrease (but still, 91% overhead)
- Size 2024 first point very computation time > overhead

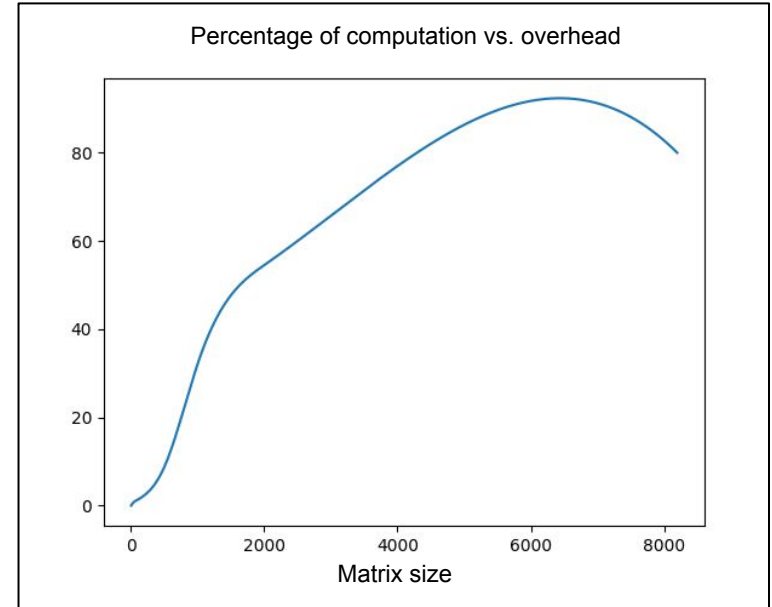


Figure: Spline (with k=3) displaying ratio computation vs. overhead.

# OpenMP - Perf and possible bottlenecks

- **No change** between actual work across the programs
- Makes sense as the program **semantics do not change** at all - only the way work is being executed
- **65% of program are arithmetics**
  - ⇒ Straightforward element-wise operations
  - ⇒ Theoretically suitable SIMD use case
  - ⇒ Compiler is already taking care of it
- **30% of program are reading, writing and moving of elements**
  - ⇒ Use matrix instead of element-wise operations
  - ⇒ Again, compiler is already taking care of it

Serial	Samples: 4K of event 'cycles', 4000 Hz, Event count (approx.): 4661322483		
	Serial::ForwardEliminationOptimized /home/marcelbraasch/CLionProjects/openmp/serialge [Percent: local period]		
	2,55	vinserf128	\$0x1,0x10(%rbx,%rax,1),%ymm7,%ymm1
	7,37	vmovupd	(%rdx,%rax,1),%xmm6
	12,54	vinserf128	\$0x1,0x10(%rdx,%rax,1),%ymm6,%ymm0
	17,44	vmulpd	%ymm3,%ymm1,%ymm1
	16,41	vsubpd	%ymm1,%ymm0,%ymm0
	2,45	vmovups	%xmm0,(%rdx,%rax,1)
	6,23	vextractf128	\$0x1,%ymm0,0x10(%rdx,%rax,1)
	31,91	add	\$0x20,%rax
	2,21	cmp	%rsi,%rax
Tasking	Samples: 13K of event 'cycles', 4000 Hz, Event count (approx.): 11385415991		
	OMP::ForwardElimination /home/marcelbraasch/CLionProjects/openmp/ompge [Percent: local period]		
	8,20	vinserf128	\$0x1,0x10(%r14,%r13,1),%ymm5,%ymm0
	12,77	vmulpd	%ymm3,%ymm1,%ymm1
	16,99	vsubpd	%ymm1,%ymm0,%ymm0
	3,67	vmovups	%xmm0,(%r14,%r13,1)
	7,45	vextractf128	\$0x1,%ymm0,0x10(%r14,%r13,1)
	33,84	add	\$0x20,%r13
	0,87	cmp	%rsi,%r13
Parallel for	Samples: 27K of event 'cycles', 4000 Hz, Event count (approx.): 25442538013		
	OMP::ForwardElimination /home/marcelbraasch/CLionProjects/openmp/ompge [Percent: local period]		
	0,05	vmovupd	(%rax,%r13,1),%xmm6
	3,21	vinserf128	\$0x1,0x10(%rax,%r13,1),%ymm6,%ymm1
	8,45	vmovupd	(%r14,%r13,1),%xmm5
	8,24	vinserf128	\$0x1,0x10(%r14,%r13,1),%ymm5,%ymm0
	13,75	vmulpd	%ymm3,%ymm1,%ymm1
	16,93	vsubpd	%ymm1,%ymm0,%ymm0
	3,82	vmovups	%xmm0,(%r14,%r13,1)
	7,52	vextractf128	\$0x1,%ymm0,0x10(%r14,%r13,1)
	32,14	add	\$0x20,%r13
	0,84	cmp	%rsi,%r13

Figure: Perf output for the respective OMP vs. serial programmes

# MPI - Parallelized Implementation and Approach

- Initial parallelization approach
  - Calculate distribution based on current status
  - Send rows from root to other ranks
  - Get calculated rows back
- Motivation of implementation
  - Easy to implement / intuitive solution
  - Organization only done in rank 0
- Bottleneck due to high amount of communication

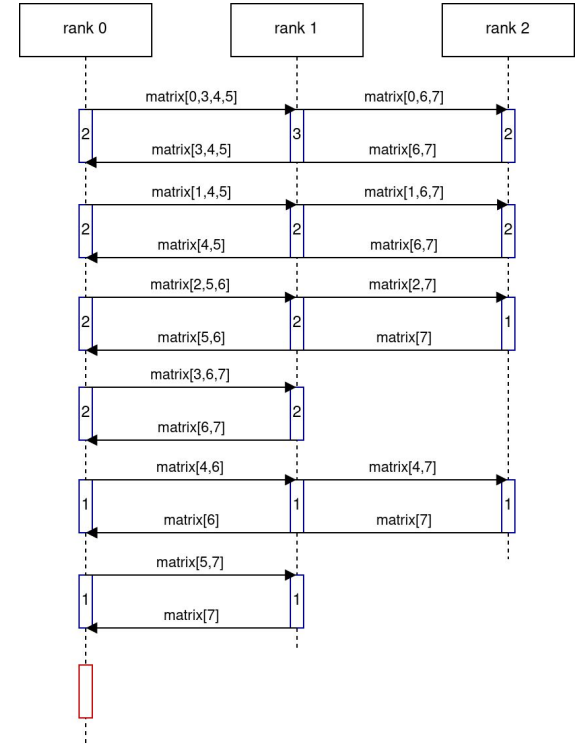


Figure: Schematic sketch of our first approach.

# MPI - Parallelized Implementation and Approach

- Bottleneck is removed by:
  - Distribute in the beginning
  - Only share current master row
- Improvements
  - Less communication overhead
  - Worse distribution of calculation load
  - More complex implementation

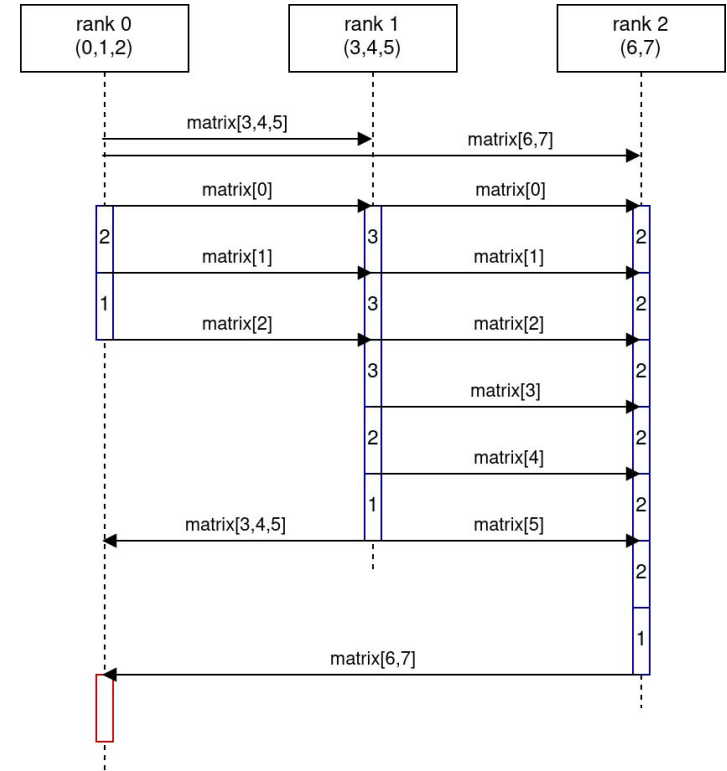
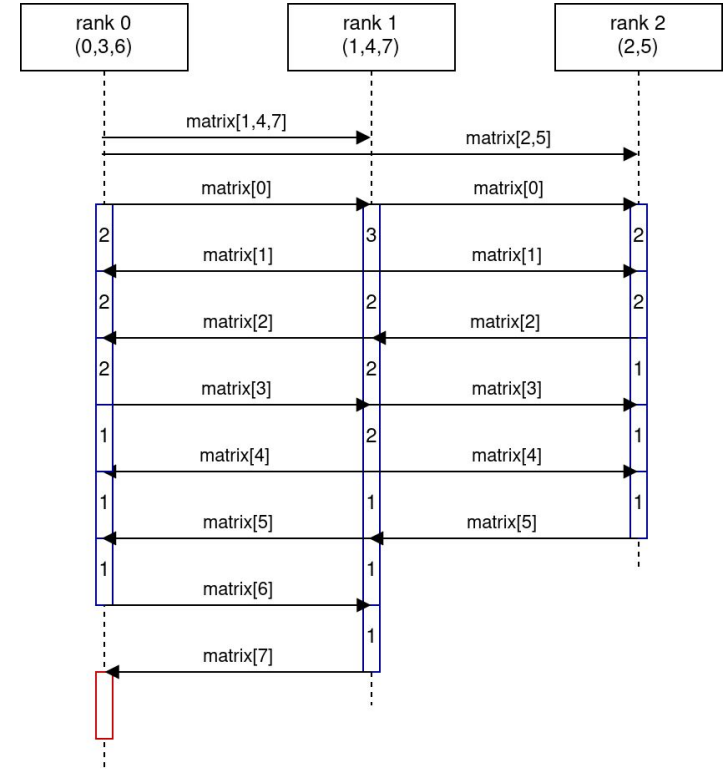


Figure: Schematic sketch of the improved approach.

# MPI - Final Implementation improvements and new speed-up

- When starting with hybrid, we developed an even better approach
  - Distribute in the beginning in round robin sequence
  - Matrix is worked off evenly by all processes



# MPI - Speed up comparison

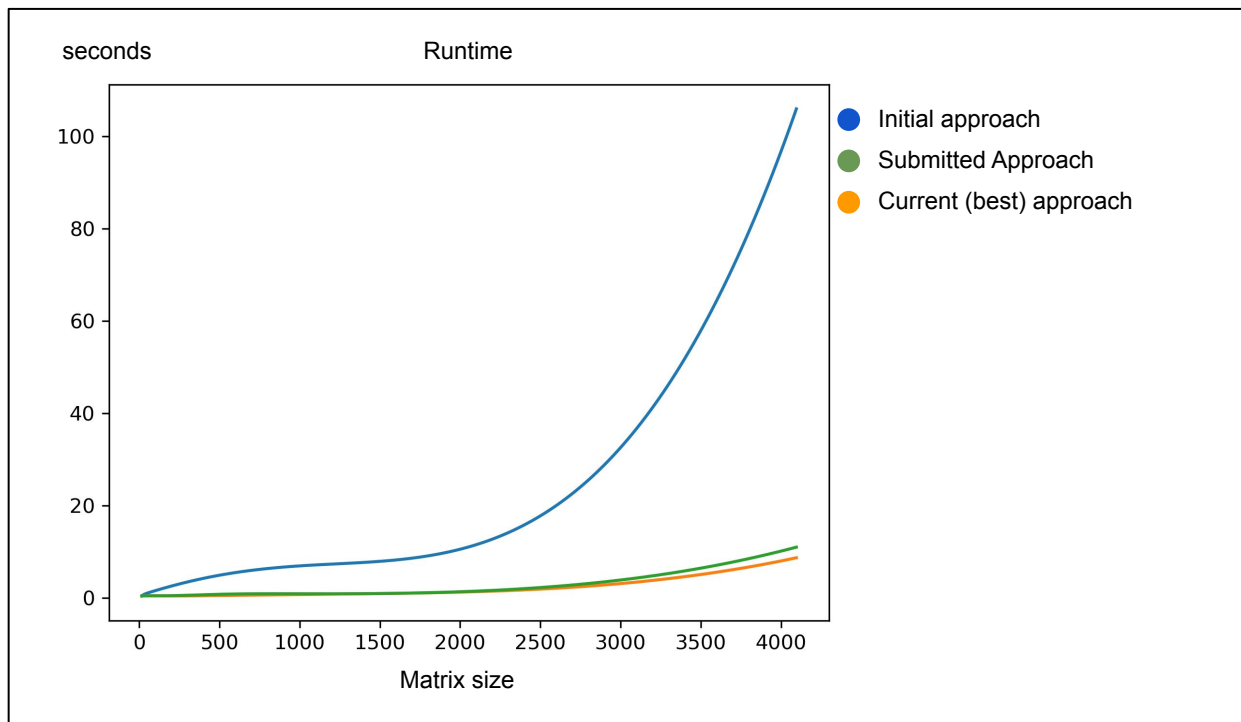


Figure: Compiler already vectorizing element-wise operations.



# MPI - Communication vs. Calculation

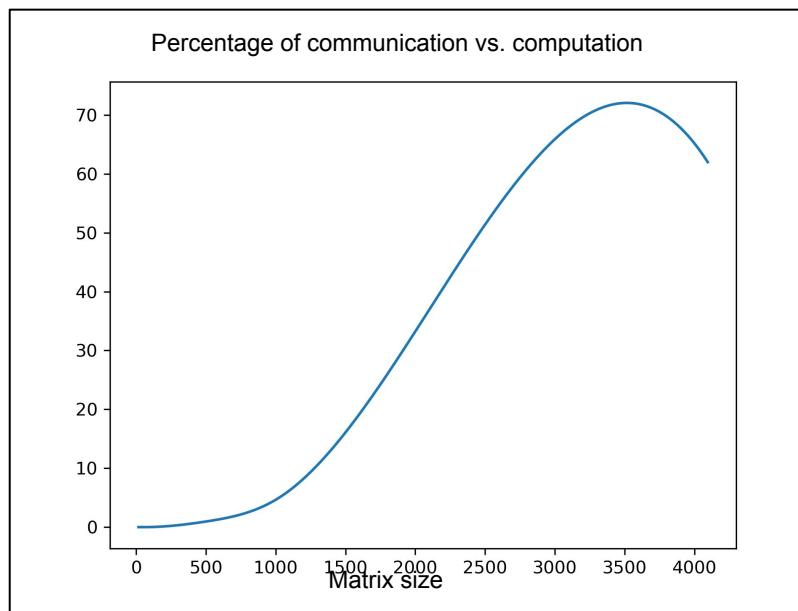
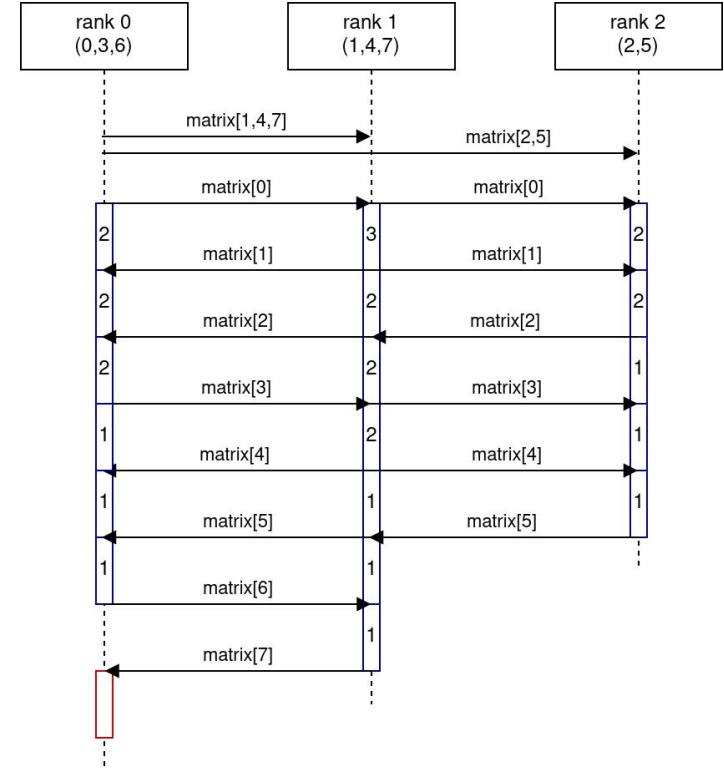


Figure: Spline (with  $k=3$ ) displaying ratio communication vs. computation using the same procedure as in the OMP part.

# Hybrid - Parallelized implementation and approach

- Built on our current best approach
- Addition of OMP similar to OMP part



# Hybrid - Final Performance Results

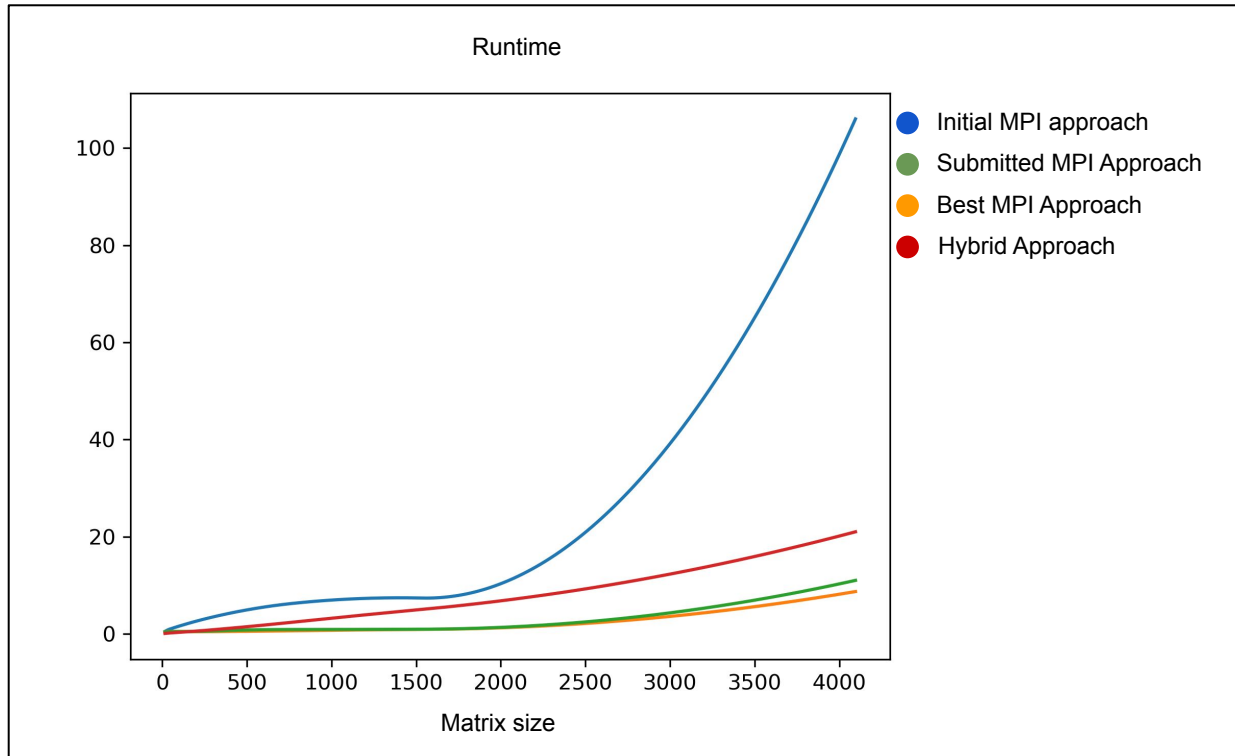


Figure: Runtime comparison between MPI

# Hybrid - Communication vs. Calculation

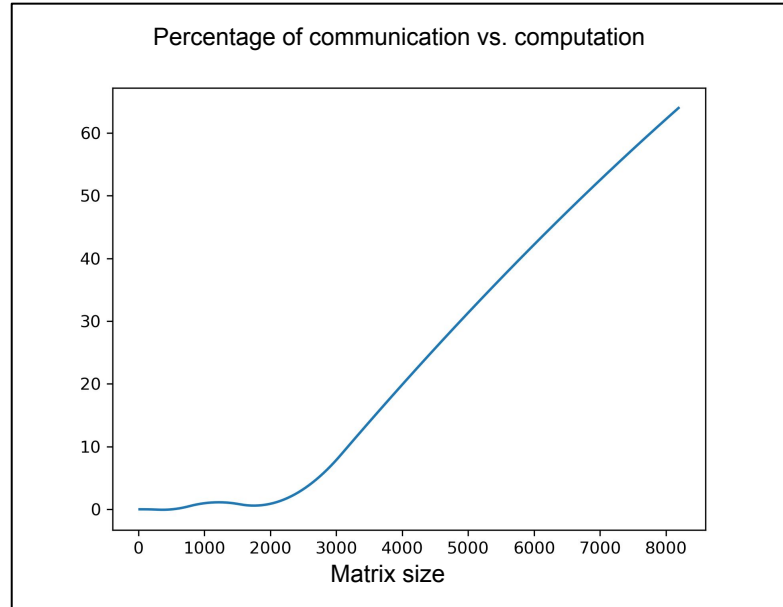


Figure: Spline (with  $k=3$ ) displaying ratio communication vs. computation using the same procedure as in the OMP part.

# Bonus - Parallelized implementation and approach

- Parallelized the inner loop using SIMD intrinsics
- Using double, we can calculate 4 values at once
- Achieved no speed-up
- Using -O3 flag compiler already figures out what to vectorize

```
Samples: 4K of event 'cycles', 4000 Hz, Event count (approx.): 4661322483
Serial::ForwardEliminationOptimized /home/marcelbraasch/CLionProjects/openmp/serialge [Percent: local period]
2,55      vinsertf128  $0x1,0x10(%rbx,%rax,1),%ymm7,%ymm1
7,37      vmovupd   (%rdx,%rax,1),%xmm6
12,54     vinsertf128 $0x1,0x10(%rdx,%rax,1),%ymm6,%ymm0
17,44     vmulpd    %ymm3,%ymm1,%ymm1
16,41     vsubpd    %ymm1,%ymm0,%ymm0
2,45      vmovups   %xmm0, (%rdx,%rax,1)
6,23      vextractf128 $0x1,%ymm0,0x10(%rdx,%rax,1)
31,91     add       $0x20,%rax
2,21      cmp       %rsi,%rax
```

Figure: Compiler already vectorizing element-wise operations.

- Using 32 threads, vectorizing rows by a factor of 4
- Theoretically, should yield a work split of  $4 \times 32 = 128$
- By Amdahl's law, theoretical max speed up is then
  - $1 / (1 - 0.97 + 0.97 / (4 \times 32)) = 26.61$

# Conclusion

- OMP outperforms all other up to 2048
- MPI is faster for matrices with  $n \geq 4092$
- Hybrid did not work very well compared to MPI
- The best we got was a speed up of 14
- Amdahl's law indicates a theoretical max of 17
- The compiler is already vectorizing matrix operations

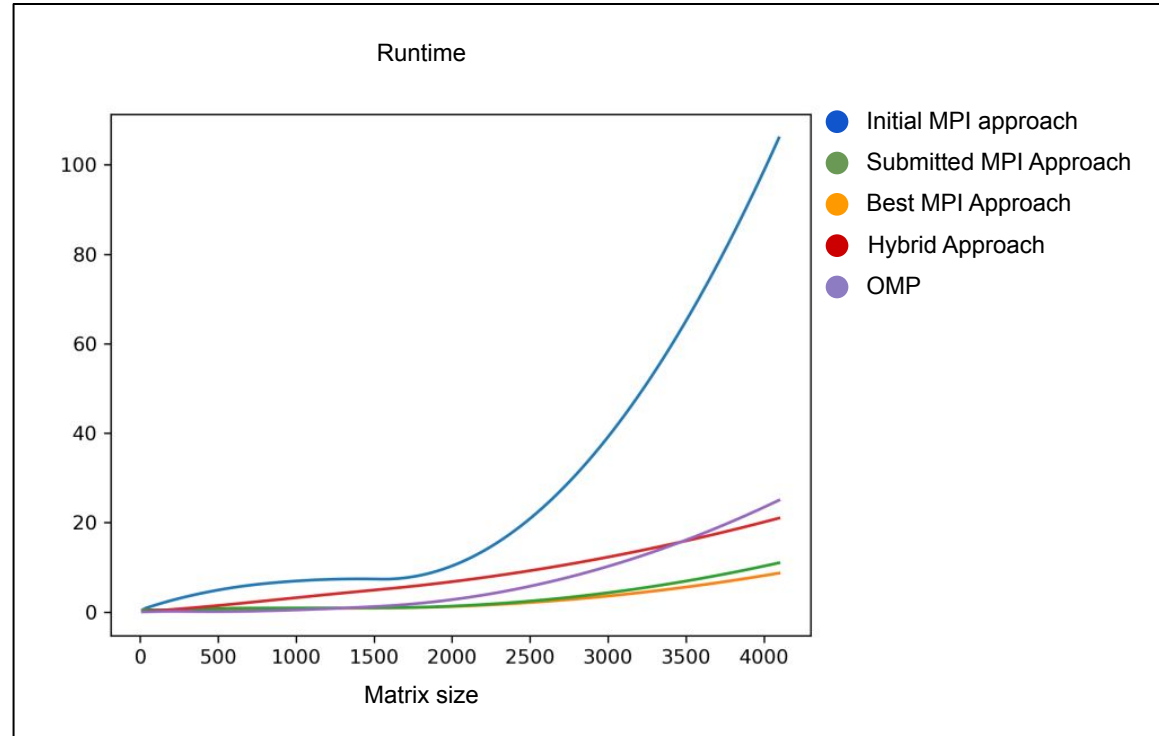


Figure: Runtime comparison between MPI

# Additional



System specifications we performed the tests on:

- OS: Ubuntu 20.04.2 LTS
- RAM: 32 GB
- CPU: AMD Ryzen 9 5950X
- GPU: Nvidia GeForce RTX 3090